



COMPUTER NETWORKS

EXP 2



Basic Functions Used for Socket Programming

JULY 30, 2021

ROEHIT RANGANATHAN
RA1911033010017 | L2

Aim:

To Discuss Some OF The Basic Functions Used For Socket Programming

Description:

1. Man socket

```
SOCKET(2)          Linux Programmer's Manual          SOCKET(2)

NAME
    socket - create an endpoint for communication

SYNOPSIS
    #include <sys/types.h>          /* See NOTES */
    #include <sys/socket.h>

    int socket(int domain, int type, int protocol);

DESCRIPTION
    socket() creates an endpoint for communication and returns a
    file descriptor that refers to that endpoint. The file
    descriptor returned by a successful call will be the lowest-
    numbered file descriptor not currently open for the process.

    The domain argument specifies a communication domain; this
    selects the protocol family which will be used for communi-
    cation. These families are defined in <sys/socket.h>. The
    currently understood formats include:

        Name                Purpose                Man page
        AF_UNIX, AF_LOCAL   Local communication    unix(7)
        AF_INET              IPv4 Internet protocols  ip(7)
        AF_INET6             IPv6 Internet protocols  ipv6(7)
        AF_IPX               IPX - Novell protocols    netlink(7)
        AF_NETLINK           Kernel user interface device
    )
        AF_X25              ITU-T X.25 / ISO-8208 protocol  x25(7)
        AF_AX25              Amateur radio AX.25 protocol
        AF_ATMPVC            Access to raw ATM PVCs
        AF_APPLETALK         AppleTalk                ddp(7)
        AF_PACKET            Low level packet interface  packet(7)
        AF_ALG               Interface to kernel crypto API
```

2. SOCK_STREAM

```
SOCK_STREAM    Provides sequenced, reliable, two-way, connection-based byte streams. An out-of-band data transmission mechanism may be supported.
```

3. SOCK_DGRAM

```
SOCK_DGRAM    Supports datagrams (connectionless, unreliable
messages of a fixed maximum length).
```

4. SOCK_SEQPACKET

```
SOCK_SEQPACKET Provides a sequenced, reliable, two-way c
onnection-based data transmission path for datagrams of fixed maximu
m length; a consumer is required to read an
entire packet with each input system call.
```

5. SOCK_RAW

```
SOCK_RAW Provides raw network protocol access.
```

6. SOCK_RDM

```
SOCK_RDM Provides a reliable datagram layer that does n
ot guarantee ordering.
```

7. SOCK_PACKET

```
SOCK_PACKET Obsolete and should not be used in new programs; see packet(7).
```

8. man connect

```
CONNECT(2) Linux Programmer's Manual CONNECT(2)

NAME
    connect - initiate a connection on a socket

SYNOPSIS
    Manual page connect(2) line 1 (press h for help or q to quit)
```

9. man accept

```

ACCEPT(2)                                Linux Programmer's Manual                                ACCEPT(2)

NAME
    accept, accept4 - accept a connection on a socket

SYNOPSIS
    #include <sys/types.h>                /* See NOTES */
    #include <sys/socket.h>

    int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);

    #define _GNU_SOURCE                    /* See feature_test_macros(7) */
    #include <sys/socket.h>

    int accept4(int sockfd, struct sockaddr *addr,
                 socklen_t *addrlen, int flags);

DESCRIPTION
    The accept() system call is used with connection-based socket types (SOCK_STREAM, SOCK_SEQPACKET). It extracts the first connection request on the queue of pending connections for the listening socket, sockfd, creates a new connected socket, and returns a new file descriptor referring to that socket. The newly created socket is not in the listening state. The original socket sockfd is unaffected by this call.

    The argument sockfd is a socket that has been created with socket(2), bound to a local address with bind(2), and is listening for connections after a listen(2).

    The argument addr is a pointer to a sockaddr structure. This structure is filled in with the address of the peer socket, as known to the communications layer. The exact format of the address returned addr is determined by the

```

10. man send

```

SEND(2)                                Linux Programmer's Manual                                SEND(2)

NAME
    send, sendto, sendmsg - send a message on a socket

SYNOPSIS
    #include <sys/types.h>
    #include <sys/socket.h>

    ssize_t send(int sockfd, const void *buf, size_t len, int flags);

    ssize_t sendto(int sockfd, const void *buf, size_t len, int flags,
                   const struct sockaddr *dest_addr, socklen_t addrlen);

    ssize_t sendmsg(int sockfd, const struct msghdr *msg, int flags);

DESCRIPTION
    The system calls send(), sendto(), and sendmsg() are used to transmit a message to another socket.

    The send() call may be used only when the socket is in a connected state (so that the intended recipient is known). The only difference between send() and write(2) is the presence of flags. With a zero flags argument, send() is equivalent to write(2). Also, the following call

        send(sockfd, buf, len, flags);

    is equivalent to

```

11. man recv

```

RECV(2)                                Linux Programmer's Manual          RECV(2)

NAME
    recv, recvfrom, recvmsg - receive a message from a socket

SYNOPSIS

```

12. man read

```

READ(2)                                Linux Programmer's Manual          READ(2)

NAME
    read - read from a file descriptor

```

13. man write

```

WRITE(1)                                BSD General Commands Manual        WRITE(1)

NAME
    write - send a message to another user

SYNOPSIS
    write user [tty]

DESCRIPTION
    The write utility allows you to communicate with other users,
    by copying lines from your terminal to theirs.

    When you run the write command, the user you are writing to
    gets a message of the form:

        Message from yourname@yourhost on yourtty at hh:mm ...

    Any further lines you enter will be copied to the specified
    user's terminal. If the other user wants to reply, they must
    run write as well.

    When you are done, type an end-of-file or interrupt character.
    The other user will see the message 'EOF' indicating that the
    conversation is over.

    You can prevent people (other than the super-user) from writ-
    ing to you with the mesg(1) command.

    If the user you want to write to is logged in on more than one
    terminal, you can specify which terminal to write to by speci-
    fying the terminal name as the second operand to the write
    command. Alternatively, you can let write select one of the
    terminals - it will pick the one with the shortest idle time.
    This is so that if the user is logged in at work and also
    dialed up from home, the message will go to the right place.

```

14. man bind

```

BIND(2)                                Linux Programmer's Manual                                BIND(2)

NAME
    bind - bind a name to a socket

SYNOPSIS
    #include <sys/types.h>           /* See NOTES */
    #include <sys/socket.h>

    int bind(int sockfd, const struct sockaddr *addr,
              socklen_t addrlen);

DESCRIPTION
    When a socket is created with socket(2), it exists in a name
    space (address family) but has no address assigned to it.
    bind() assigns the address specified by addr to the socket
    referred to by the file descriptor sockfd. addrlen speci-
    fies the size, in bytes, of the address structure pointed to
    by addr. Traditionally, this operation is called "assigning
    a name to a socket".

    It is normally necessary to assign a local address using
    bind() before a SOCK_STREAM socket may receive connections
    (see accept(2)).

    The rules used in name binding vary between address fami-
    lies. Consult the manual entries in Section 7 for detailed
    information. For AF_INET, see ip(7); for AF_INET6, see
    ipv6(7); for AF_UNIX, see unix(7); for AF_APPLETALK, see
    ddp(7); for AF_PACKET, see packet(7); for AF_X25, see
    x25(7); and for AF_NETLINK, see netlink(7).

    The actual structure passed for the addr argument will
    depend on the address family. The sockaddr structure is
    defined as something like:

```

15. ifconfig

```

Kkottilingam:~/environment $ ifconfig
Kkottilingam:~/environment $ ifconfig
docker0: flags=4099<UP,BROADCAST,MULTICAST> mtu 1500
    inet 172.17.0.1 netmask 255.255.0.0 broadcast 172.17.255.25
    5
    ether 02:42:ef:25:cc:24 txqueuelen 0 (Ethernet)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 9001
    inet 172.31.9.200 netmask 255.255.240.0 broadcast 172.31.15
    .255
    inet6 fe80::47b:6fff:fe64:1e5d prefixlen 64 scopeid 0x20<li
    nk>
    ether 06:7b:6f:64:1e:5d txqueuelen 1000 (Ethernet)
    RX packets 332007 bytes 244600134 (244.6 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 195655 bytes 49685414 (49.6 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 666 bytes 95406 (95.4 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 666 bytes 95406 (95.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

16. man htons/man htonl

```

BYTEORDER(3)          Linux Programmer's Manual          BYTEORDER(3)

NAME
    htonl, htons, ntohs, ntohl - convert values between host and
    network byte order

SYNOPSIS
    #include <arpa/inet.h>

    uint32_t htonl(uint32_t hostlong);

    uint16_t htons(uint16_t hostshort);

    uint32_t ntohl(uint32_t netlong);

    uint16_t ntohs(uint16_t netshort);

DESCRIPTION
    The htonl() function converts the unsigned integer hostlong
    from host byte order to network byte order.

    The htons() function converts the unsigned short integer
hostshort from host byte order to network byte order.

    The ntohl() function converts the unsigned integer netlong
    from network byte order to host byte order.

    The ntohs() function converts the unsigned short integer
netshort from network byte order to host byte order.

    On the i386 the host byte order is Least Significant Byte
    first, whereas the network byte order, as used on the Inter-
    net, is Most Significant Byte first.

ATTRIBUTES
    For an explanation of the terms used in this section, see
    attributes(7).

```

17. man gethostname

```

GETHOSTNAME(2)        Linux Programmer's Manual        GETHOSTNAME(2)

NAME
    gethostname, sethostname - get/set hostname

SYNOPSIS
    #include <unistd.h>

    int gethostname(char *name, size_t len);
    int sethostname(const char *name, size_t len);

Feature Test Macro Requirements for glibc (see fea-
ture_test_macros(7)):

    gethostname():
        Since glibc 2.12: _BSD_SOURCE || _XOPEN_SOURCE >= 500
        || /* Since glibc 2.12: */ _POSIX_C_SOURCE >= 200112L
    sethostname():
        Since glibc 2.21:
            _DEFAULT_SOURCE
        In glibc 2.19 and 2.20:
            _DEFAULT_SOURCE || (_XOPEN_SOURCE && _XOPEN_SOURCE < 5
00)

        Up to and including glibc 2.19:
            _BSD_SOURCE || (_XOPEN_SOURCE && _XOPEN_SOURCE < 500)

DESCRIPTION
    These system calls are used to access or to change the host-
    name of the current processor.

    sethostname() sets the hostname to the value given in the
    character array name. The len argument specifies the number
    of bytes in name. (Thus, name does not require a terminat-
    ing null byte.)

```

18. gethostbyname

```
GETHOSTBYNAME(3)      Linux Programmer's Manual      GETHOSTBYNAME(3)

NAME
    gethostbyname, gethostbyaddr, sethostent, gethostent, end-
    hostent, h_errno, perror, hstrerror, gethostbyaddr_r, get-
    hostbyname2, gethostbyname2_r, gethostbyname_r, gethostent_r
    - get network host entry

SYNOPSIS
    #include <netdb.h>
    extern int h_errno;

    struct hostent *gethostbyname(const char *name);

    #include <sys/socket.h>      /* for AF_INET */
    struct hostent *gethostbyaddr(const void *addr,
                                   socklen_t len, int type);

    void sethostent(int stayopen);

    void endhostent(void);

    void perror(const char *s);

    const char *hstrerror(int err);

    /* System V/POSIX extension */
    struct hostent *gethostent(void);

    /* GNU extensions */
    struct hostent *gethostbyname2(const char *name, int af);

    int gethostent_r(
        struct hostent *ret, char *buf, size_t buflen,
        struct hostent **result, int *h_errnop);
```